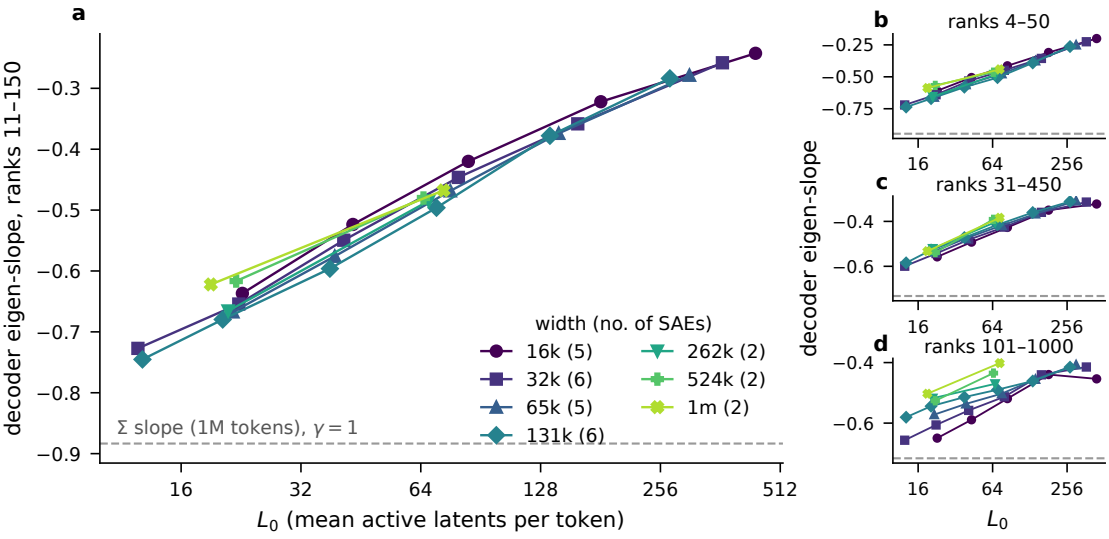
a
b
c
d
ranks 4–50
ranks 31–450
ranks 101–1000
decoder eigen-slope, ranks 11–150
decoder eigen-slope
Σ slope (1M tokens), γ = 1
width (no. of SAEs)
16k (5)
32k (6)
65k (5)
131k (6)
262k (2)
524k (2)
1m (2)
L_0 (mean active latents per token)
L_0