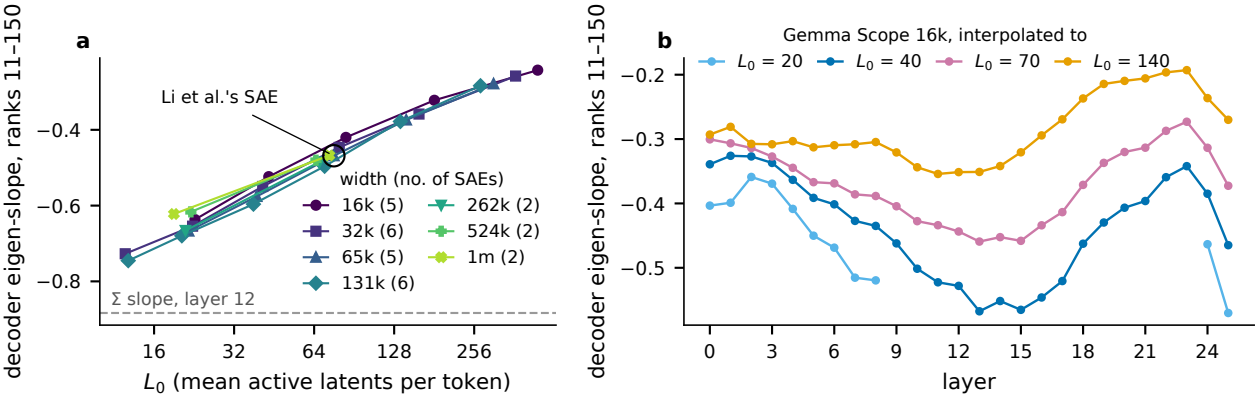
a
Li et al.'s SAE
width (no. of SAEs)
16k (5)
32k (6)
65k (5)
131k (6)
262k (2)
524k (2)
1m (2)
Σ slope, layer 12
decoder eigen-slope, ranks 11–150
−0.4
−0.6
−0.8
16
32
64
128
256
L_0 (mean active latents per token)
b
Gemma Scope 16k, interpolated to
L_0 = 20
L_0 = 40
L_0 = 70
L_0 = 140
decoder eigen-slope, ranks 11–150
−0.2
−0.3
−0.4
−0.5
0
3
6
9
12
15
18
21
24
layer