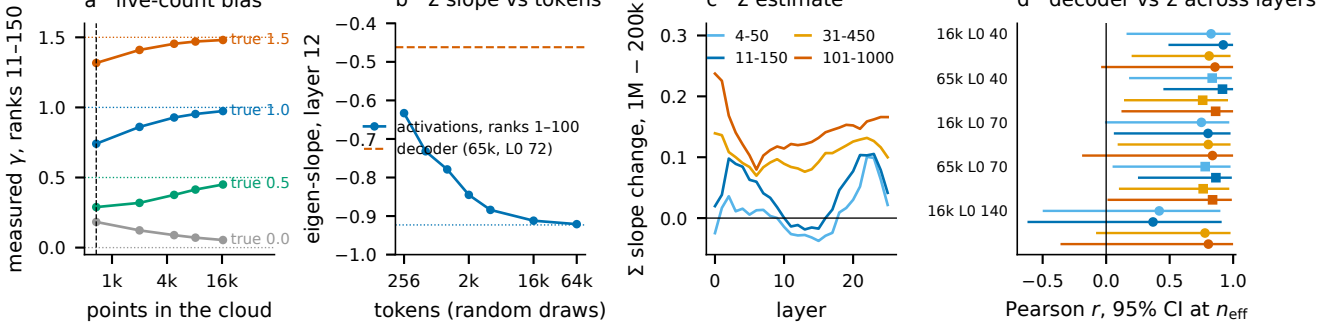
a live-count bias
measured γ, ranks 11–150
1.5
1.0
0.5
0.0
true 1.5
true 1.0
true 0.5
true 0.0
1k
4k
16k
points in the cloud
b Σ slope vs tokens
eigen-slope, layer 12
−0.4
−0.5
−0.6
−0.7
−0.8
−0.9
−1.0
activations, ranks 1–100
decoder (65k, L0 72)
256
2k
16k
64k
tokens (random draws)
c Σ estimate
Σ slope change, 1M − 200k
0.3
0.2
0.1
0.0
4-50
11-150
31-450
101-1000
0
10
20
layer
d decoder vs Σ across layers
16k L0 40
65k L0 40
16k L0 70
65k L0 70
16k L0 140
−0.5
0.0
0.5
1.0
Pearson r, 95% CI at n_{eff}